

# CS 486



# Personal Notes

---

Marcus Chan

Taught by Pascal Poupart & Sriram Ganapathi  
Subramanian

JW CS '25



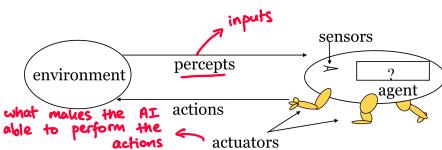
# Chapter 1: Introduction

## REINFORCEMENT LEARNING PROBLEM



Our goal is for the AI to learn to choose actions that maximize rewards.

## AGENTS & ENVIRONMENTS



The "agent function" maps percepts to actions; ie  $f: P \rightarrow A$ .

The "agent program" runs on the physical architecture to produce  $f$ .

## RATIONAL AGENTS

A "rational agent" chooses whichever action that maximizes the expected value of its performance measure given the percept sequence to date.

Note that rationality is not omniscience, but rather learning & autonomy.

## PEAS

"PEAS" helps us specify the task environment:

- ① Performance measure;  
eg safety, destination, etc
- ② Environment;  
eg streets, traffic, etc
- ③ Actuators; &  
eg steering, brakes, etc.
- ④ Sensors.  
eg GPS, engine sensors, etc.

## PROPERTIES OF TASK ENVIRONMENTS

Task environments can be:

- ① fully vs partially observable:
  - fully observable: agent knows state of the world from the stimuli
  - partially observable: agent does not directly observe the world's state

② deterministic vs stochastic;

- deterministic: next state is observable at any time
- stochastic: next state is unpredictable

③ episodic vs sequential;

- episodic: agent's current action will not affect a future action
- sequential: agent's current action will affect a future action

④ static vs dynamic;

- static: model is trained once
- dynamic: model is trained continuously

⑤ discrete vs continuous;

⑥ single agent vs multiagent.

\*the former option is "easier" than the latter.

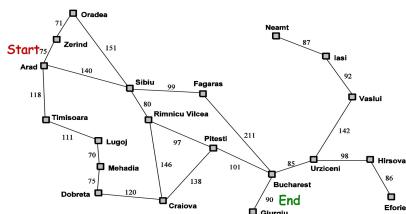
# Chapter 2: Uninformed Search Techniques

## SIMPLE PROBLEM SOLVING AGENT

```
function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  static: seq, an action sequence, initially empty
    state, some description of the current world state
    goal, a goal, initially null
    problem, a problem formulation
  state ← UPDATE-STATE(state, percept)
  if seq is empty then do
    goal ← FORMULATE-GOAL(state)
    problem ← FORMULATE-PROBLEM(state, goal)
    seq ← SEARCH(problem)
  action ← FIRST(seq)
  seq ← REST(seq)
  return action
```

- 💡 This can only tackle problems that are
  - ① fully observable;
  - ② deterministic;
  - ③ sequential;
  - ④ static;
  - ⑤ discrete; &
  - ⑥ single agent.

## EXAMPLE: TRAVELLING IN ROMANIA



- initial state: In Arad
- actions: drive between cities
- goal test: In Bucharest
- path cost: distance between cities

## EXAMPLE: 8-TILE PUZZLE



Start State



Goal State

- states: locations of 8 tiles & blank
- initial state: any state
- actions: up, down, left, right
- goal test: does state match desired configuration
- path cost: # of steps

## SEARCHING

- 💡 We can visualize a state space search in terms of trees or graphs:
  - ① nodes correspond to states; &
  - ② edges correspond to taking actions.
- 💡 These "search trees" are formed using "search nodes", which have
  - ① the state associated with it;
  - ② parent node & operator applied to the parent to reach the "current" node;
  - ③ cost of the path so far; &
  - ④ depth of the node.

## EXPANDING NODES

- 💡 "Expanding a node" refers to applying all legal operators to the state contained in the node & generating nodes for all corresponding successor states.



## GENERIC SEARCH ALGORITHM

- 💡 Algorithm:
  - ① Initialize search with initial problem state.
  - ② Then repeat:
    - if no candidate nodes can be expanded, return failure
    - otherwise, choose a leaf node for expansion according to our search strategy.
    - if the node contains a goal state, return the solution.
    - otherwise, expand the node by applying the legal operators to the state associated within the node, & add the resulting nodes to the tree.

# EVALUATING SEARCH ALGORITHMS

We can use the following properties when evaluating search algorithms:

- ① "completeness" — is the algorithm guaranteed to find a solution (if it exists)?
  - ② "optimality" — does the algorithm find the optimal solution (ie lowest path cost)?
  - ③ time & space complexity.
- We consider the following variables:
- ① "branching factor" ( $b$ ) — the # of children each node has
  - ② depth of shallowest goal node ( $d$ ); &
  - ③ max length of any path in the state space ( $m$ ).

## BREADTH-FIRST SEARCH

Refer to CS341 notes for details; we expand all nodes on a given level before any node on the next level is expanded.

Evaluating the algorithm:

- ① Completeness: yes if  $b \geq 0$
- ② Optimality: yes if all costs are same
- ③ Time:  $1 + b + \dots + b^d \in O(b^d)$
- ④ Space:  $O(b^d)$ .

\* all uninformed search methods have exponential time complexity.

## UNIFORM COST SEARCH

Idea: we expand the node with the lowest path cost.

We can implement this using a priority queue.

Let  $C^*$  = cost of optimal solution &  $\epsilon = \min \text{action cost}$ . Then

- ① Completeness: yes if  $\epsilon > 0$
- ② Optimality: yes
- ③ Time:  $O(b^{C^*/\epsilon})$
- ④ Space:  $O(b^{C^*/\epsilon})$

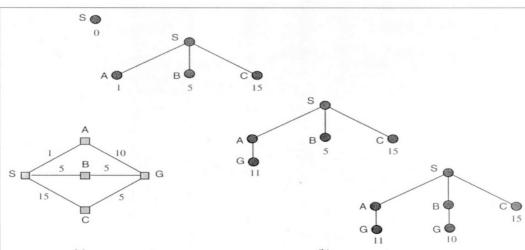


Figure 3.13 A route-finding problem. (a) The state space, showing the cost for each operator. (b) Progression of the search. Each node is labelled with  $g(n)$ . At the next step, the goal node with  $g = 10$  will be selected.

## DEPTH-FIRST SEARCH

Refer to CS341 notes for details; we expand the deepest node in the current fringe of the search tree first.

Evaluation:

- ① Complete: no  
- may get stuck going down a long path
- ② Optimal: no  
- might return a solution which is deeper (ie more costly) than another solution
- ③ Time:  $O(b^m)$   
- note we might have  $m > d$
- ④ Space:  $O(bm)$

## DEPTH-LIMITED SEARCH

Idea: Treat all nodes at depth  $l$  as if they have no successors.  
- try to choose  $l$  based on the problem

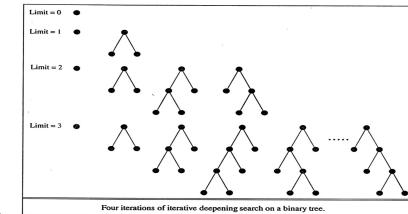
This avoids the problem of unbounded trees.

Evaluation:

- ① Time:  $O(b^l)$
- ② Space:  $O(b^l)$
- ③ Complete: no
- ④ Optimal: no

## ITERATIVE-DEEPENING

Idea: repeatedly perform depth-limited search, but increase the limit each time.



Evaluation:

- ① Complete: yes
- ② Optimal: yes
- ③ Time:  $O(b^d)$
- ④ Space:  $O(b^d)$

Time:

$$\begin{aligned}
 \text{(limit=1)} & \quad 1 \\
 \text{(limit=2)} & \quad 1 + b \\
 \text{(limit=3)} & \quad 1 + b + b^2 \\
 & \vdots \\
 \text{(+) (limit=d)} & \quad 1 + b + b^2 + \dots + b^d \\
 & \quad d + (d-1)b + (d-2)b^2 + \dots \\
 & \quad + b^d \in O(b^d)
 \end{aligned}$$

# Chapter 3: Informed Search Techniques

## MOTIVATION

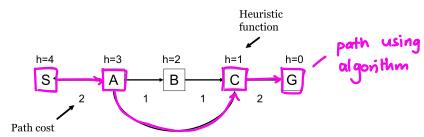
- Q<sub>1</sub>: In search problems, we often have additional knowledge about the problem.  
eg with the "travelling around Romania", we know dist. bw cities  
↳ so we can find the overhead in going the wrong direction
- Q<sub>2</sub>: Our knowledge is often about the "merit" of nodes.
- Q<sub>3</sub>: Notions of merit:
  - ① how expensive it is to get from a state to a goal;
  - ② how easy it is to get from a state to a goal.

## HEURISTIC FUNCTIONS • $h(n)$

- Q<sub>1</sub>: We need to develop domain specific "heuristic functions"  $h(n)$ , which guess the cost of reaching the goal from node  $n$ .
- Q<sub>2</sub>: In general, if  $h(n_1) < h(n_2)$ , we guess reaching the goal is cheaper from  $n_1$  than from  $n_2$ . We also need
  - ①  $h(n) = 0$  if  $n$  is a goal node
  - ②  $h(n) > 0$  if  $n$  is not a goal node.

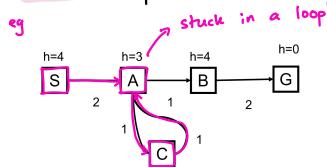
## GREEDY BEST-FIRST SEARCH

- Q<sub>1</sub>: Idea: Use  $h(n)$  to rank the nodes in the fringe & expand the node with the lowest  $h$ -value.  
(ie "greedily" trying to find the least-cost solution).
- Q<sub>2</sub>: Example:



- Q<sub>3</sub>: Note greedy best-first is not optimal.  
eg in the above example,
  - path found has cost=6.
  - but cheaper path is  $S \rightarrow A \rightarrow B \rightarrow C \rightarrow G$  with cost=6.

- Q<sub>4</sub>: It is also not complete, as it can be stuck in loops.



- eg - but if we check for repeated states then we are okay
- Q<sub>5</sub>: This algorithm uses exponential space & worst-case time.

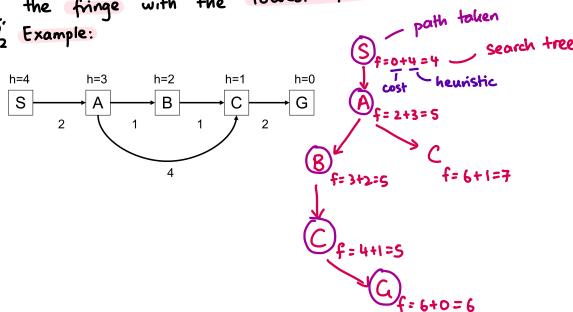
## A\* SEARCH

Idea: Define  $f(n) = g(n) + h(n)$ , where

- ①  $g(n)$  = cost of path to node  $n$
- ②  $h(n)$  = heuristic estimate of cost of reaching goal from node  $n$ .

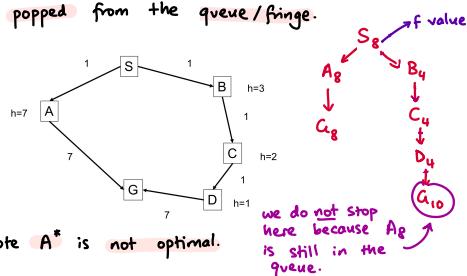
We then iteratively expand the node in the fringe with the lowest  $f$  value.

Example:

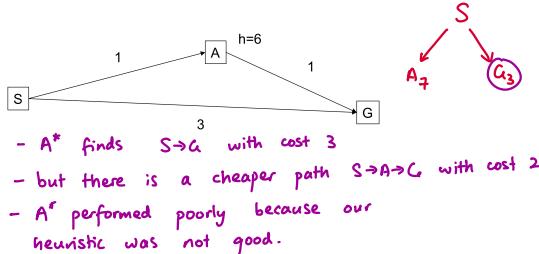


$A^*$  should terminate only when the goal state is popped from the queue/fringe.

eg



Note  $A^*$  is not optimal.



- A\* finds  $S \rightarrow A$  with cost 3
- but there is a cheaper path  $S \rightarrow A \rightarrow G$  with cost 2.
- A\* performed poorly because our heuristic was not good.

## ADMISSIBLE [HEURISTICS]

Let  $h^*(n)$  be the true minimal cost to the goal from node  $n$ .

Then, we say the heuristic  $h(n)$  is "admissible" if

$$h(n) \leq h^*(n) \quad \forall n.$$

In particular, admissible heuristics never overestimate the cost to the goal.

$h(n)$  IS ADMISSIBLE  $\Rightarrow A^*$  IS OPTIMAL

If  $h(n)$  is admissible, then  $A^*$  with tree-search is optimal.

Proof: let  $G$  be an optimal goal state, &  $f(G) = f^* = g(G)$ .

Let  $G_2$  be a suboptimal goal state, ie  $f(G_2) = g(G_2) > f^*$ . (since  $h(G_2) = h(G_2) = 0$ )

Assume, for a cont'n, that  $A^*$  selects  $G_2$  from the queues; ie  $A^*$  terminates with a suboptimal solution.

Let  $n =$  node currently a leaf node on an optimal path to  $G$ .



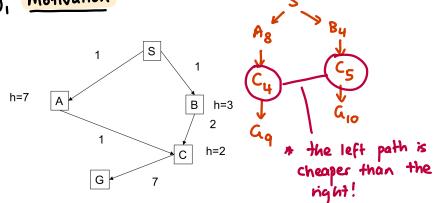
$G_2$

$\bullet G_2$

Since  $h$  is admissible,  $f^* \geq f(n)$ . If  $n$  is not chosen for expansion over  $G_2$ , then  $f(n) \geq f(G_2)$ . Thus  $f^* \geq f(G_2)$ . Since  $h(G_2) = 0$ , thus  $f^* \geq g(G_2)$ , a cont'n.

## REVISITING STATES IN A\*

Motivation:



If we allow states to be expanded again, we might get a better solution!

## CONSISTENT [HEURISTICS]

We say  $h(n)$  is "consistent" if

$$h(n) \leq \text{cost}(n, n') + h(n') \quad \forall n, n'.$$

Note that  $A^*$  graph-search with a consistent heuristic is optimal.

## PROPERTIES OF A\*

Note that  $A^*$  is

- ① Complete if  $h(n)$  is consistent; -  $f$  always increases along any path
- ② Has exponential time complexity in the worst-case; & - but a good heuristic helps a lot -  $O(\text{cbm})$  if heuristic is perfect
- ③ Has exponential space complexity.

## ITERATIVE DEEPENING A\*

(IDA\*)

- Idea: Like iterative deepening search, but change f-cost rather than depth in each iteration.
- This reduces the space complexity.

## SIMPLIFIED MEMORY-BOUNDED A\*

(SMA\*)

- Idea: Proceeds like A\* but when it runs out of memory it drops the worst leaf node (ie one with highest f-value).
- If all leaf nodes have the same f-value, then drop the oldest & expand the newest.
- This is
  - ① optimal;
  - ② complete if depth of shallowest goal node < memory size.

## OBTAINING HEURISTICS

- One approach to get heuristics is to think of an easier problem & let  $h(n)$  be the cost of reaching the goal in the easier problem.

We can also

- ① precompute solution costs of subproblems & store them in a pattern database; or
- ② learn from experience with the problem class.

## EXAMPLE: 8-PUZZLE GAME

We can relax the game in 3 ways:

- ① We can move tile from position A  $\rightarrow$  B if A is next to B (ignore whether position is blank)
- ② We can move tile from position A  $\rightarrow$  B if B is blank (ignore adjacency)
- ③ We can move tile from position A  $\rightarrow$  B regardless.

③ leads to the "misplaced tile heuristic". ( $h_3$ )

- to solve this problem we need to move each tile into its final position.
- # of moves = # of misplaced tiles
- admissible

① leads to the "manhattan distance heuristic". ( $h_1$ )

- to solve this we need to slide each tile into its final position

- admissible

Note  $h_1$  "dominates"  $h_3$ ; ie  $h_3(n) \leq h_1(n) \forall n$ .

# Chapter 4:

# Constraint Satisfaction

## INTRODUCTION

Q: These are useful for problems where the state structure is important.

Q: In many problems, the same state can be reached independent of the order in which the moves are chosen.

Q: So, we can try to solve problems efficiently by being smart about the action order.

## 4-QUEENS CONSTRAINT PROPAGATION

Q: Idea: Remove conflicting squares from consideration when we put a queen down.



## CONSTRAINT SATISFACTION PROBLEM (CSP)

Q: A "constraint satisfaction problem" is defined by some  $\{V, D, C\}$ , where

①  $V = \{V_1, \dots, V_n\}$  is a set of variables;

②  $D = \{D_1, \dots, D_n\}$  is a set of domains, where  $D_i$  is the set of possible values for each  $V_i$ ;

&

③  $C = \{C_1, \dots, C_m\}$  is the set of constraints.

## STATE

Q: A "state" is an assignment of values to some or all of the variables;

ie  $V_i = x_i, V_j = x_j, \dots$ .

## CONSISTENT [ASSIGNMENT]

Q: We say an assignment is "consistent" if it does not violate any constraints.

## SOLUTION

Q: A "solution" is a complete, consistent assignment.

## EXAMPLE: 8 QUEENS AS A CSP

Q: 8 queens as a CSP:

- variables:  $V_{ij}, i, j = 1, \dots, 8$

- domain of each var:  $\{0, 1\}$

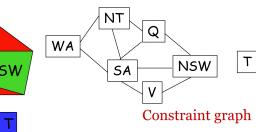
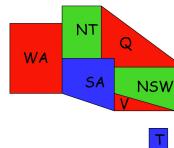
- constraints:  $V_{ij} = 1 \Rightarrow V_{ik} = 0 \quad \forall k \neq j$

$V_{ij} = 1 \Rightarrow V_{kj} = 0 \quad \forall k \neq i$

similar constraint for diagonals

$$\sum_{i,j} V_{ij} = 8$$

## EXAMPLE: MAP COLORING



Constraint graph

MAP COLORING AS A CSP

- variables: WA, NT, ..., T (the regions)

- each var has the same domain: {red, green, blue}

- no 2 adjacent variables have the same value

(ie  $WA \neq NT$ ,  $WA \neq SA$ , etc)

## PROPERTIES OF CSPS

**Q:** Types of variables:

① Discrete & finite;

eg 8-queens, map coloring

\* we focus on this in this course.

② Discrete with infinite domains; &

eg job scheduling

③ Continuous domains.

**Q:** Types of constraints:

① "Unary constraint": relates a single variable to a value

- eg Queensland = blue

② "Binary constraint": relates two variables

③ "Higher order constraints": relates  $\geq 3$  variables.

## CSPs & SEARCH

**Q:** We can formulate CSPs as a search problem:

① we have  $N$  variables  $V_1, \dots, V_n$ ;

② a valid assignment is  $\{V_i = x_i, \dots, V_n = x_n\}$ ,  $0 \leq i \leq n$  where values satisfy the variable constraints.

③ states: valid assignments

④ initial state: empty assignment

⑤ successor:  $\{V_i = x_i, \dots, V_n = x_n\} \rightarrow \{V_i = x_i, \dots, V_n = x_n, V_{n+1} = x_{n+1}\}$

⑥ goal test: complete assignment

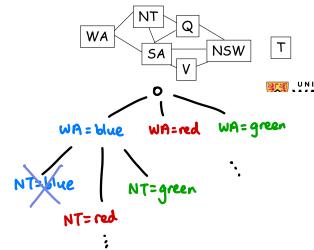
## BACKTRACKING SEARCH

```

function BACKTRACKING-SEARCH(csp) returns a solution, or failure
    return RECURSIVE-BACKTRACKING({}, csp)
function RECURSIVE-BACKTRACKING(assignment, csp) returns a solution, or failure
    if assignment is complete then return assignment
    var  $\leftarrow$  SELECT-UNASSIGNED-VARIABLE(Variables[csp], assignment, csp)
    for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
        if value is consistent with assignment according to Constraints[csp] then
            add { var = value } to assignment
            result  $\leftarrow$  RECURSIVE-BACKTRACKING(assignment, csp)
            if result  $\neq$  failure then return result
            remove { var = value } from assignment
    return failure
  
```

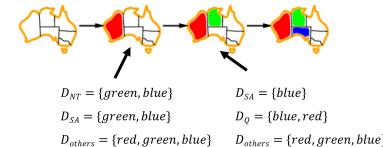
- this is DFS that choose values for one variable at a time
- we "backtrack" when a variable has no legal values to assign

## EXAMPLE: MAP COLORING



## MOST CONSTRAINED VARIABLE HEURISTIC

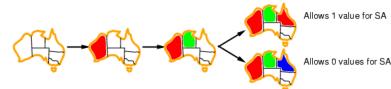
**Q:** Idea: Choose the variable which has the fewest "legal" moves.



**Q:** In a tie, choose the variable with the most constraints on the remaining variables.  
(ie "most constraining variable").

## LEAST CONSTRAINING VALUE HEURISTIC

**Q:** Idea: Given a variable, choose the "least constraining value", ie the one that rules out the fewest values in the remaining variables.

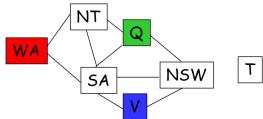


## FORWARD CHECKING

Q<sub>1</sub>: Idea: We keep track of remaining legal values for unassigned variables, & terminate search when any variable has no legal values.

Q<sub>2</sub>: This helps us detect failure early.

### EXAMPLE: MAP COLORING



WA	NT	Q	NSW	V	SA	T
RGB						
R	GB	RGB	RGB	RGB	GB	RGB
R	B	G	RB	RGB	B	RGB
R	B	G	RB	B	X	RGB

WA = R  
Q = G  
V = B

this is the empty set;  
⇒ the current assignment does not lead to a solution.

# Chapter 5: Uncertainty

Q<sub>1</sub>: Refer to STAT 231/330 for more details.

Q<sub>2</sub>: We use  $\sim$  to denote the complement of an event (ie  $\sim A$ ).

## BAYES RULE

Q<sub>1</sub>: For 2 events A, B, note

$$P(B|A) = \frac{P(A|B)P(B)}{P(A)}$$

Proof:  $P(A)P(B|A) = P(A \wedge B) = P(B)P(A|B)$ .

$$\therefore P(B|A) = \frac{P(B)P(A|B)}{P(A)}.$$

Q<sub>2</sub>: In particular, it allows us to compute a belief about hypothesis B given evidence A.

Q<sub>3</sub>: More general forms:

$$P(A|B) = \frac{P(B|A)P(A)}{P(B|A)P(A) + P(B|\sim A)P(\sim A)}$$

$$P(A|B \wedge X) = \frac{P(B|A \wedge X)P(A|X)}{P(B|X)}$$

$$P(A=v_i|B) = \frac{P(B|A=v_i)P(A=v_i)}{\sum_{k=1}^n P(B|A=v_k)P(A=v_k)}$$

## PROBABILISTIC INFERENCE

Q<sub>1</sub>: Idea: Given a prior distribution  $P(X)$  over variables  $X$  of interest & given new evidence  $E=e$  for some variable E, revise our degrees of belief; ie the "posterior"  $P(X|E=e)$ .

## ISSUES

Q<sub>1</sub>: Specifying the full joint distribution for  $X_1, \dots, X_n$  requires an exponential number of possible "worlds".

Q<sub>2</sub>: So, inference is also slow since we need to sum over these exponential number of worlds:

$$P(X_1, \dots, X_{i-1}, X_{i+1}, \dots, X_n | X_i) = \frac{P(X_1, \dots, X_n)}{\sum_{x_1} \sum_{x_{i-1}} \sum_{x_{i+1}} \dots \sum_{x_n} P(X_1, \dots, X_n)}$$

## CONDITIONAL INDEPENDENCE

Q<sub>1</sub>: Two variables  $X, Y$  are "conditionally independent" given  $Z$  if

$$P(X=x | Z=z) = P(X=x | Y=y, Z=z)$$

$$\Leftrightarrow P(X=x, Y=y | Z=z) = P(X=x | Z=z)P(Y=y | Z=z)$$

$$\Leftrightarrow \forall x \in \text{dom}(X), y \in \text{dom}(Y), z \in \text{dom}(Z)$$

Q<sub>2</sub>: If we know the value of  $Z$ , nothing we learn about  $Y$  will influence our beliefs about  $X$ .

## VALUE OF INDEPENDENCE

Q<sub>1</sub>: If  $X_1, \dots, X_n$  are mutually independent, then we can specify the full joint distribution using only  $n$  parameters (ie linear) instead of  $2^n - 1$  (ie exponential).

Q<sub>2</sub>: Although most domains do not exhibit complete mutual independence, they do instead exhibit a fair amount of conditional independence.

## NOTATION: $P(X)$

Q<sub>1</sub>: We define " $P(X)$ " as the marginal distribution over  $X$ .

-  $P(X=x)$  is a number,  $P(X)$  is a distribution.

## NOTATION: $P(X|Y)$

Q<sub>1</sub>: We define " $P(X|Y)$ " as the family of conditional distributions over  $X$ ; one for each  $y \in \text{dom}(Y)$ .

# EXPLOITING CONDITIONAL INDEPENDENCE: CHAIN RULE

Consider a story:

- If Pascal woke up too early  $E$ , Pascal probably needs coffee  $C$ ; if Pascal needs coffee, he's likely grumpy  $G$ . If he is grumpy then it's possible that the lecture won't go smoothly  $L$ . If the lecture does not go smoothly then the students will likely be sad  $S$ .



$E$  - Pascal woke up too early     $G$  - Pascal is grumpy     $S$  - Students are sad  
 $C$  - Pascal needs coffee     $L$  - The lecture did not go smoothly

$S$  is independent of  $E, C, G$  given  $L$   
 $L$  is independent of  $E, C$  given  $G$  & so on.

$$\Rightarrow P(S|L, G, C, E) = P(S|L)$$

$$P(L|G, C, E) = P(L|G)$$

$$P(G|C, E) = P(G|C)$$

Then

$$\begin{aligned} P(S, L, G, C, E) &= P(S|L, G, C, E) P(L|G, C, E) P(G|C, E) \cdot \\ &\quad P(C|E) P(E) \\ &= \underline{P(S|L) P(L|G) P(G|C) P(C|E) P(E)}. \end{aligned}$$

Q: In this, we can specify the full joint distribution by specifying the five local conditional distributions.